

How *Degeneracy* Helps in Streaming Triangle Counting: A Detailed Review

Main Results and Contributions: Bera & Seshadhri (2020) study triangle counting in graph streams when the input graph has *bounded degeneracy* (denoted κ). Their **main positive result** is a constant-pass streaming algorithm that uses space $\widetilde{O}(m\kappa/T)$ to $(1 \pm \epsilon)$ -approximate T , the number of triangles, in an m -edge input graph of degeneracy $\leq \kappa$. In particular, for constant- κ graph families (e.g. planar or minor-closed families, or preferential-attachment social graphs), this space is $\widetilde{O}(m/T)$, which is asymptotically *smaller* than the previous worst-case bounds of $\min(m^{3/2}/T, m/\sqrt{T})$ for arbitrary graphs. The **main lower bound** (Theorem 1.3) nearly matches this: any constant-pass algorithm must use $\Omega(m\kappa/T)$ space. Thus their results identify the precise $(m\kappa/T)$ dependence, up to polylog factors. This closes the gap between worst-case streaming bounds and the degeneracy-sensitive case. For example, on an n -vertex *wheel graph* (a sparse graph of degeneracy $\kappa=3$ with $O(n)$ edges and $\Theta(n)$ triangles), the new bound is essentially logarithmic space, whereas prior algorithms would require $\Omega(n)$ or $\Omega(\sqrt{n})$ space.

- **Streaming Algorithm (Theorem 1.2):** A constant-pass insertion-stream algorithm returns a $(1 \pm \epsilon)$ -approximation to the triangle count using $\widetilde{O}(m\kappa/T)$ space. This leverages the graph's low degeneracy; intuitively, edges can be directed and sampled according to a low "edge-degeneracy" (minimum endpoint degree) ordering.
- **Lower Bound (Theorem 1.3):** Any (multi-pass) streaming algorithm must use $\Omega(m\kappa/T)$ space, even allowing deletions. This nearly matches the upper bound and shows the degeneracy dependence is necessary.
- **Significance:** For bounded- κ families, the space is $\widetilde{O}(m/T)$, a dramatic improvement. The wheel-graph example shows that for $\kappa=O(1)$ graphs, the algorithm's space can be polylogarithmic in n , whereas prior generic bounds were polynomial. In summary, the paper introduces the first streaming triangle-count algorithm whose complexity explicitly improves when the graph has low degeneracy, breaking the "triangle-deletion" lower bounds for such graphs.

Secondary/Supporting Results: In support of the main theorems, the paper presents a careful algorithmic framework (Algorithms 1–2) using *edge-oriented sampling* and rigorous variance analysis. Key ideas include: directing edges by degeneracy ordering (à la Chiba–Nishizeki) and assigning each triangle to its least "heavy" edge to control overlaps. Lemmas establish that the estimator is unbiased and concentrated, and quantify the space used. In addition, the authors provide illustrative examples (like the wheel graph) and comparisons with prior work, emphasizing how constant degeneracy graphs form a broad class (including planar, bounded-expansion, and preferential-attachment graphs ¹) where their bound is strongest. They also show that existing lower bounds (e.g. by Bera–Chakrabarti (2017), McGregor–Vorochnikova (2016)) are subsumed in the general $\Omega(m\kappa/T)$ tradeoff ² ³.

Future Research Directions & Open Questions: The paper highlights several intriguing directions. Most broadly, it suggests examining *other* streaming graph problems under low-degeneracy assumptions. For example, known successes include graph coloring ⁴, matching size estimation ⁵, and independent-set approximation ⁶ that exploit sparsity. An explicit open question is whether *fixed-size subgraph counting*

(beyond triangles) admits faster streaming algorithms on low- κ graphs. In fact, the authors **conjecture** that for counting r -cliques, a similar $\tilde{O}(m\kappa/T)$ space bound should hold:

Conjecture 7.1: “Consider a graph with degeneracy κ that has many r -cliques. There exists a constant-pass streaming algorithm that $(1 \pm \epsilon)$ -approximates the number of r -cliques using $\tilde{O}(m\kappa/T)$ bits of space.” ⁷ .

They pose this to generalize their triangle result to larger cliques. More generally, the authors ask: can any other sublinear-time or streaming problems exploit bounded degeneracy to beat worst-case bounds? (E.g. estimating the count of other patterns.) Finally, they note that all known streaming lower bounds for triangle counting “break” on low- κ inputs, so proving tight bounds or creating new algorithms in that regime is fertile ground.

Real-world Applications (in use today): While the paper is theoretical, it addresses a practical scenario. Many real-world networks have *low degeneracy*. As the authors note, bounded-degeneracy graphs include planar graphs, minor-closed families, and models of social networks like preferential-attachment graphs ¹ . In practice, triangle counts in such networks underpin metrics like clustering coefficient and community structure. Any streaming analysis of massive social or web graphs could benefit: for a network with constant κ , the space becomes $\tilde{O}(m/T)$, enabling triangle estimation on graphs too large to fit in memory. In principle, this algorithm could be deployed in streaming analytics systems (e.g. real-time social graph monitoring, network security for detecting dense subnets, or large-scale biological network motif counting) where edges arrive rapidly. For example, in monitoring a social network feed, one might want on-the-fly triangle statistics in the limited memory of a streaming engine; this result shows that if the network is sparse enough (low degeneracy), only a tiny sample of edges is needed. (Indeed, the authors emphasize “low degeneracy is an often observed characteristic of real-world graphs” ⁸ , supporting the practical relevance.) Moreover, foundational work on graph processing (e.g. Chiba–Nishizeki’s arboricity-based listing algorithms) has shown that degeneracy-orientations speed up subgraph tasks, and this streaming result effectively portends such optimizations to the streaming model.

Potential Future Applications: Looking forward, these methods could be used in any domain requiring counting of small substructures in massive graphs under memory constraints. For example, in **network motif discovery** (biology), one often counts triangles or cliques in protein-interaction networks, which tend to be sparse. A degeneracy-aware streaming tool would allow faster approximate motif counts. In **graph databases** and query engines, join-size estimation often reduces to triangle counting; exploiting degeneracy could yield faster query planning. In **social-media analytics** (e.g. counting closed triads in real-time message graphs for community detection), the improved space bounds could allow more granular monitoring. Looking even further, if Conjecture 7.1 (and related extensions) holds, we might see streaming algorithms for larger clique or subgraph counts in sparse graphs – for instance, real-time detection of 4-cliques in communication networks to signal emergent communities. In short, any streaming graph analytics system stands to gain if the graph is known to have bounded degeneracy.

Citing Works (selected): We found the following significant works that explicitly cite or build on Bera–Seshadhri (2020):

- “Approximately Counting Subgraphs in Data Streams” – Hendrik Fichtenberger and Pan Peng (2022, *arXiv:2203.14225*). This paper develops general streaming algorithms to count arbitrary constant-size subgraphs in insertion-only streams. Notably, it gives a constant-pass algorithm for

approximating the number of r -cliques in a graph of degeneracy λ using $O(m\lambda^{r-2}/T)$ space ⁹. In doing so, it *resolves a conjecture by Bera and Seshadhri* regarding clique-counting in bounded-degeneracy graphs ¹⁰. In particular, Fichtenberger & Peng improve the space for cliques of size r by showing that degeneracy λ yields an $O(m\lambda^{r-2}/T)$ bound, and they explicitly reference the Bera–Seshadhri PODS 2020 result as motivation ¹⁰. Thus this work extends the triangle-counting ideas to general subgraphs, using the degeneracy framework developed in Bera–Seshadhri. (It appears as arXiv preprint 2022 and later in a conference.)

(We found no other direct citations in the literature indexed by Google Scholar, Semantic Scholar, DBLP, or arXiv beyond this. Many recent works on subgraph counting and streaming leverage related ideas of graph sparsity and degeneracy, but do not explicitly cite Bera–Seshadhri 2020. For completeness, related works include Bera et al. (2020, ITCS) on subgraph counting in sparse graphs and Pashanasangi–Seshadhri (2021, KDD) on temporal triangle counting via degeneracy ordering, but these treat offline or temporal settings and do not directly cite the PODS 2020 paper.)

Sources: The above summary is based on Bera & Seshadhri (PODS 2020) and its text ⁸ ⁷, together with Fichtenberger & Peng (arXiv 2022) ¹⁰. All numerical results and quotes are cited verbatim from these sources.

¹ ² ³ ⁴ ⁵ ⁶ ⁷ ⁸ [2003.13151] How the Degeneracy Helps for Triangle Counting in Graph Streams

<https://ar5iv.org/pdf/2003.13151>

⁹ ¹⁰ arxiv.org

<https://arxiv.org/pdf/2203.14225>